

IBM Bob 2.0 hackathon · team Trust Me Bro

TrustMeBro

Your agent keeps the roadmap.

Plan-first projects for AI coding agents: map the code, put every new task on the roadmap before touching it, keep the records current, and let the next session start where the last one stopped.

Multi-session projects with an AI agent lose their thread.

Planning is manual. Turning an idea into phases, tasks, "done" criteria and risks takes days, and the plan ends up in scattered notes.

Every session starts cold. The developer re-explains where things stand, or the whole history is dumped into the context window.

We lived it

Our first IBM Bob planning session used
~160,000 tokens without writing a line of code,
~80,000 of them from fetching the same page twice.

Three files the agent keeps current.

Map it	<code>/layout-init</code> writes <code>LAYOUT.md</code> : folders, components, database, where each piece of logic lives.
Plan it	<code>/roadmap-planner</code> turns an idea into <code>PLAN.md</code> , waits for approval, then creates <code>ROADMAP.md</code> .
Record first	Every new task goes on the roadmap before the code: feature, fix, or big work behind an approved plan.
Tick with evidence	Items are ticked with the file, test or result that proves them; the layout is updated too.
Resume cold	The next session starts from the progress headline and next actions, not from rereading the code.

Copy `.bob/` into a project. Bob does the rest.

Rules in `.bob/rules/trustmebro.md`, always loaded.

Hooks: `SessionStart` loads progress and next actions; `UserPromptSubmit` creates `ROADMAP.md` and reminds Bob: record first.

Skills: layout-init, roadmap-planner, roadmap-sync.

Live Bob Shell run: "Add multiply(a, b)"

Todo list step 1: add the task to `ROADMAP.md`.
Roadmap entry **before** the code, then
`LAYOUT.md` updated, item ticked with evidence,
progress refreshed to 100%.

35 s · 12 tool calls · all checks passed

Same prompt on a half-finished project.

"Continue the project: do the next open task." New session, same model, same code: shoplite, 35 files, 16 tests, 3 features done, 2 open, 1 known bug.

Base agent

Explores the code to work out what is done and next, writes the change, records nothing.

104 s · \$0.30 · 593k tokens · 27 tool calls

With TrustMeBro

Reads the next action from the roadmap, does it, ticks it in `ROADMAP.md`.

51 s · \$0.22 · 468k tokens · 13 tool calls

Three setups, four scenarios.

Scenario	Setup	Correct	Time	Tokens	Cost	Roadmap current
Resume the project	Base	3/3	104 s	593k	\$0.303	n/a
	TrustMeBro	3/3	51 s	468k	\$0.215	3/3
What's done, what's left?	Base	2 of 3	23 s	116k	\$0.083	n/a
	TrustMeBro	3 of 3	11 s	63k	\$0.076	n/a
Add a cross-module feature	Base	3/3	58 s	368k	\$0.201	n/a
	Docs only	3/3	57 s	406k	\$0.210	0/3
	TrustMeBro	3/3	69 s	485k	\$0.266	3/3
Two tasks, two sessions	Base	2/2	114 s	966k	\$0.407	n/a
	TrustMeBro	2/2	149 s	1,020k	\$0.520	2/2

What it shows

Faster to pick up. Honest about the cost.

2× faster

resuming a project, 29% cheaper

46% fewer tokens

answering "what's left?", and it found the bug
only the roadmap knew

100%

of change runs kept the roadmap current; docs
without the plugin: 0 of 5

+28–32%

cost on pure implementation tasks, spent on
keeping the roadmap and layout up to date

Same quality

every setup did every change task correctly

Works with

One repo, six agents.

IBM Bob tested

copy `.bob/` into the project

Claude Code tested

```
/plugin marketplace add  
TonyLikeDev/tricklord
```

GitHub Copilot CLI tested

```
copilot plugin install  
trustmebro@trustmebro
```

Antigravity CLI validated

```
agy plugin install ./tricklord/src
```

Codex untested

```
codex plugin add trustmebro@trustmebro
```

Gemini CLI untested

```
gemini extensions install ./tricklord/src
```

Thank you

Your agent keeps the roadmap.

github.com/TonyLikeDev/tricklord

Website in `web/` · benchmark in `benchmark/` · Bob test in `trustmebro-bob-test/`